

Colab link:<https://colab.research.google.com/drive/1GpsP-3l95M0hyZ8Dv8ptrjBJPuxGyj1n>
(<https://colab.research.google.com/drive/1GpsP-3l95M0hyZ8Dv8ptrjBJPuxGyj1n>)

Colab link models:

<https://colab.research.google.com/drive/1JpXZ9XXCwny86wpeltn0psLIJtw1aD7K#scrollTo=ddtLILDmJ7ng>
(<https://colab.research.google.com/drive/1JpXZ9XXCwny86wpeltn0psLIJtw1aD7K#scrollTo=ddtLILDmJ7ng>)

Introduction

This project will employ machine learning to classify different newsgroups data and explore additional trends and observations. The dataset used is the well-known 20 Newsgroups dataset, which has become quite frequency used in the field of machine learning. The 20 Newsgroups dataset is a collection of approximately 20,000 newsgroup documents, partitioned (nearly) evenly across 20 different newsgroups. It was originally collected by Ken Lang, probably for his Newsweeder: Learning to filter netnews paper, though he does not explicitly mention this collection. The reasoning behind choosing newsgroups data is mainly connected to also having real life applications. Through an advanced neural network it is possible to classify a certain piece of news simply by its content. This can be applied when uncertainty occurs about what category certain news might belong to. In this case the journalist can use the algorithm to help determine what specific category it might belong to. Additionally the algorithm can be used to suggest another category which can make the news fall into multiple categories. This can save time from various journalists/news publishers, who are no longer required to manually determine what category certain news belong to. This project will start by exploring the data with the use of a Glove embedding model, where the link between different words is being researched. Afterwards a network analysis will be constructed to explore how pair of words are connect with the use of Bigrams. After exploring the text, a baseline Random Forest model will be constructed to see how well it performs. Subsequently a neural network will be constructed with the goal of classifying multiple topics. The neural network used in this project is "Long Short Term Memory", which is a superior neural network for sequential text data. Lastly the project will conclude upon the discoveries made.

Data preparation

First, we read in the data and inspect it. The inspection has the purpose of checking that the data is as we expect, for this we use the glimpse command and find nothing unexpected. Next, we check the topic variable to both see if there are the 20 categories we expect, and to find if there is a somewhat even distribution between the 20 topics.

```
In [1]: install.packages("tm")
install.packages("data.table")
install.packages("text2vec")
install.packages("quanteda")
install.packages("tidytext")
install.packages("textstem")
install.packages("dplyr")
install.packages("wider")
install.packages("ggforce") # Awesome plotting
install.packages("ggraph")
install.packages("uwot")
install.packages("igraph") # For network analysis
install.packages("ggraph")
install.packages("tidygraph")
install.packages("mapplots")
install.packages("tidyr")
library(tidyr)
library(tidyverse)
library(data.table)
library(tm)
library(mapplots)
library(ggforce)
library(ggraph)
library(uwot)
library(igraph)
library(ggraph)
library(tidygraph)
library(wider)
library(dplyr)
library(textstem)
library(text2vec)
library(quanteda)
library(tidytext)

data=read_csv("https://www.dropbox.com/s/y29q10pyrkdtz67/Newspaper.csv?dl=1")
```

Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
also installing the dependencies 'NLP', 'slam'

Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
also installing the dependencies 'formatR', 'iterators', 'lambda.r', 'futile.
options', 'RcppParallel', 'foreach', 'irlba', 'futile.logger', 'mlapi', 'spar
sepp'

Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
also installing the dependencies 'coda', 'extrafontdb', 'Rttf2pt1', 'RcppEige
n', 'statnet.common', 'reticulate', 'ISOcodes', 'extrafont', 'fastmatch', 'gg
repel', 'network', 'RSpectra', 'sna', 'SnowballC', 'spacyr', 'stopwords', 'pr
oxyC', 'RcppArmadillo'

Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
also installing the dependencies 'hunspell', 'tokenizers', 'jane austenr'

Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
also installing the dependencies 'zoo', 'dtt', 'syllly.en', 'syllly', 'syuzhe
t', 'english', 'mgsub', 'qdapRegex', 'koRpus.lang.en', 'koRpus', 'lexicon',
'textclean', 'textshape'

Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
also installing the dependencies 'farver', 'tweenr', 'polyclip'

Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
also installing the dependencies 'gridExtra', 'igraph', 'viridis', 'tidygrap
h', 'graphlayouts'

Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
also installing the dependencies 'sitmo', 'FNN', 'RcppAnnoy', 'RcppProgress',
'dqrng'

Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)
Installing package into '/usr/local/lib/R/site-library'
(as 'lib' is unspecified)

— **Attaching packages** — tidyverse 1.2.1

```
✓ ggplot2 3.2.1    ✓ purrr  0.3.3
✓ tibble  2.1.3    ✓ dplyr  0.8.3
✓ readr   1.3.1    ✓ stringr 1.4.0
✓ ggplot2 3.2.1    ✓ forcats 0.4.0
```

— **Conflicts** — tidyverse_conflicts()

```
✗ dplyr::filter() masks stats::filter()
✗ dplyr::lag()     masks stats::lag()
```

Attaching package: 'data.table'

The following objects are masked from 'package:dplyr':

between, first, last

The following object is masked from 'package:purrr':

transpose

Loading required package: NLP

Attaching package: 'NLP'

The following object is masked from 'package:ggplot2':

annotate

Loading required package: Matrix

Attaching package: 'Matrix'

The following objects are masked from 'package:tidyr':

expand, pack, unpack

Attaching package: 'igraph'

The following objects are masked from 'package:dplyr':

as_data_frame, groups, union

The following objects are masked from 'package:purrr':

compose, simplify

The following object is masked from 'package:tibble':

as_data_frame

The following object is masked from 'package:tidyr':

crossing

The following objects are masked from 'package:stats':

decompose, spectrum

The following object is masked from 'package:base':

union

Attaching package: 'tidygraph'

The following object is masked from 'package:igraph':

groups

The following object is masked from 'package:stats':

filter

Loading required package: koRpus.lang.en

Loading required package: koRpus

Loading required package: sylly

For information on available language packages for 'koRpus', run

available.koRpus.lang()

and see ?install.koRpus.lang()

Attaching package: 'koRpus'

The following object is masked from 'package:readr':

tokenize

Attaching package: 'text2vec'

The following object is masked from 'package:igraph':

normalize

Package version: 1.5.2

Parallel computing: 2 of 2 threads used.

See <https://quanteda.io> for tutorials and examples.

Attaching package: 'quanteda'

The following objects are masked from 'package:koRpus':

tokens, types

The following objects are masked from 'package:tidygraph':

as.igraph, convert

The following object is masked from 'package:igraph':

as.igraph

The following objects are masked from 'package:tm':

as.DocumentTermMatrix, stopwords

The following object is masked from 'jupyter:irkernel':

View

The following object is masked from 'package:utils':

View

Warning message:

"Missing column names filled in: 'X1' [1]"Parsed with column specification:

```
cols(
  X1 = col_double(),
  text_id = col_double(),
  text = col_character(),
  category = col_character()
)
```

In [0]: data %<>% mutate(topic=category) %>% select(-category)

In [3]: data %>% glimpse

Observations: 18,773

Variables: 4

```
$ X1      <dbl> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 1
8...
$ text_id <dbl> 1366, 6368, 7292, 3951, 7399, 5424, 2652, 2680, 5195, 4420,
3...
$ text    <chr> "From: csfed@uxa.ecn.bgu.edu (Frank Doss)\r\n Subject: Re: S
c...
$ topic   <chr> "['alt.atheism']", "['sci.space']", "['talk.religion.misc']",
...
```

```
In [4]: data %>% head(1)
```

A tibble: 1 × 4

X1	text_id	text	topic
<dbl>	<dbl>	<chr>	<chr>
		From: csfed@uxa.ecn.bgu.edu (Frank Doss) Subject: Re: Science and theories Organization: Educational Computing Network Lines: 27 NNTP-Posting-Host: uxa.ecn.bgu.edu	
		In article <C5u7Bq.J43@news.cso.uiuc.edu> cobb@alexia.lis.uiuc.edu (Mike Cobb) writes:	
		>The examples he gave were quarks and continental plates. Are there	
		Sounds like more of the same. Gods were used to describe almost everything in the past. As we come to understand the underpinnings of more and more, the less we credit to a god. Now, the not-so-well understood elements (at least by the author) includes quarks and tectonic drift. I guess that's better than describing the perceived patterns of stars in the sky as heroes being immortalized by the gods.	
1	1366	Kinda sounds like old-earth creation--It seems that life did, indeed, evolve from a common ancestor. What caused that initial common ancestor?	[alt.atheism]
		Are we going to hear another debate on causeless events? ;-)	
		>explanations of science or parts of theories that are not measurable in and of	
		>themselves, or can everything be quantified, measured, tested, etc.?	
		>	
		>MAC	
		Michael A. Cobb	
		--	
		Frank Doss	
		The above stated words are my opinions and do not reflect the opinions, attitudes, or policies of my employer or any affiliated organizations.	

In [5]: `data %>% group_by(topic) %>% count()`

A grouped_df: 40 × 2

topic	n
<chr>	<int>
<hr/>	
['alt.atheism']	317
['comp.graphics']	389
['comp.os.ms-windows.misc']	394
['comp.sys.ibm.pc.hardware']	389
['comp.sys.mac.hardware']	381
['comp.windows.x']	394
['misc.forsale']	388
['rec.autos']	393
['rec.motorcycles']	398
['rec.sport.baseball']	396
['rec.sport.hockey']	393
['sci.crypt']	396
['sci.electronics']	390
['sci.med']	395
['sci.space']	391
['soc.religion.christian']	398
['talk.politics.guns']	364
['talk.politics.mideast']	376
['talk.politics.misc']	310
['talk.religion.misc']	251
alt.atheism	475
comp.graphics	579
comp.os.ms-windows.misc	591
comp.sys.ibm.pc.hardware	588
comp.sys.mac.hardware	567
comp.windows.x	590
misc.forsale	583
rec.autos	593
rec.motorcycles	598
rec.sport.baseball	591
rec.sport.hockey	600
sci.crypt	594
sci.electronics	588
sci.med	592

topic	n
<chr>	<int>
sci.space	593
soc.religion.christian	599
talk.politics.guns	546
talk.politics.mideast	564
talk.politics.misc	465
talk.religion.misc	374

Looking at the topic variable a glaring issue jumps out, there are duplicates of each category an issue we will address twice as we are using both R and python in our project. In each case we will redefine the variable so the categories are uniform, leaving us with a dataset that can be used in our different models further preprocessing will be required before the different models.

```
In [0]: data$topic=ifelse(data$topic=="['alt.atheism']", "atheism", data$topic)
data$topic=ifelse(data$topic=="alt.atheism", "atheism", data$topic)

data$topic=ifelse(data$topic=="['comp.graphics']", "compgraphics", data$topic)
data$topic=ifelse(data$topic=="comp.graphics", "compgraphics", data$topic)

data$topic=ifelse(data$topic=="['comp.os.ms-windows.misc']", "comp-windows", data$topic)
data$topic=ifelse(data$topic=="comp.os.ms-windows.misc", "comp-windows", data$topic)

data$topic=ifelse(data$topic=="['comp.sys.ibm.pc.hardware']", "ibm-pc-hardware", data$topic)
data$topic=ifelse(data$topic=="comp.sys.ibm.pc.hardware", "ibm-pc-hardware", data$topic)

data$topic=ifelse(data$topic=="['comp.sys.mac.hardware']", "mac-hardware", data$topic)
data$topic=ifelse(data$topic=="comp.sys.mac.hardware", "mac-hardware", data$topic)

data$topic=ifelse(data$topic=="['comp.windows.x']", "windows-x", data$topic)
data$topic=ifelse(data$topic=="comp.windows.x", "windows-x", data$topic)

data$topic=ifelse(data$topic=="['misc.forsale']", "misc-forsale", data$topic)
data$topic=ifelse(data$topic=="misc.forsale", "misc-forsale", data$topic)

data$topic=ifelse(data$topic=="['rec.autos']", "autos", data$topic)
data$topic=ifelse(data$topic=="rec.autos", "autos", data$topic)

data$topic=ifelse(data$topic=="['rec.motorcycles']", "motorcycles", data$topic)
data$topic=ifelse(data$topic=="rec.motorcycles", "motorcycles", data$topic)

data$topic=ifelse(data$topic=="['rec.sport.baseball']", "baseball", data$topic)
data$topic=ifelse(data$topic=="rec.sport.baseball", "baseball", data$topic)

data$topic=ifelse(data$topic=="['rec.sport.hockey']", "hocey", data$topic)
data$topic=ifelse(data$topic=="rec.sport.hockey", "hocey", data$topic)

data$topic=ifelse(data$topic=="['sci.crypt']", "crypt", data$topic)
data$topic=ifelse(data$topic=="sci.crypt", "crypt", data$topic)

data$topic=ifelse(data$topic=="['sci.electronics']", "electronics", data$topic)
data$topic=ifelse(data$topic=="sci.electronics", "electronics", data$topic)

data$topic=ifelse(data$topic=="['sci.med']", "medicin", data$topic)
data$topic=ifelse(data$topic=="sci.med", "medicin", data$topic)

data$topic=ifelse(data$topic=="['sci.space']", "space", data$topic)
data$topic=ifelse(data$topic=="sci.space", "space", data$topic)

data$topic=ifelse(data$topic=="['soc.religion.christian']", "christianity", data$topic)
data$topic=ifelse(data$topic=="soc.religion.christian", "christianity", data$topic)
```

```

data$topic=ifelse(data$topic=="['talk.politics.guns']", "guns", data$topic)
data$topic=ifelse(data$topic=="talk.politics.guns", "guns", data$topic)

data$topic=ifelse(data$topic=="['talk.politics.mideast']", "mideast", data$topic)
data$topic=ifelse(data$topic=="talk.politics.mideast", "mideast", data$topic)

data$topic=ifelse(data$topic=="['talk.politics.misc']", "politics", data$topic)
data$topic=ifelse(data$topic=="talk.politics.misc", "politics", data$topic)

data$topic=ifelse(data$topic=="['talk.religion.misc']", "religion-misc", data$topic)
data$topic=ifelse(data$topic=="talk.religion.misc", "religion-misc", data$topic)

```

In [7]: data %>% group_by(topic) %>% count()

A grouped_df: 20 × 2

topic	n
<chr>	<int>
atheism	792
autos	986
baseball	987
christianity	997
comp-windows	985
compgraphics	968
crypt	990
electronics	978
guns	910
hocey	993
ibm-pc-hardware	977
mac-hardware	948
medicin	987
mideast	940
misc-forsale	971
motorcycles	996
politics	775
religion-misc	625
space	984
windows-x	984

GLOVE

GloVe stands for Global Vector for Word Representation, and is an unsupervised learning algorithm, which we use for obtaining vector representations. The model rests on a simple idea that ratios of word-word co-occurrence probabilities have the potential for encoding some form of meaning which can be encoded as vector differences.

```
In [0]: data_coor <- data %>% corpus(docid_field = "topic", text_field = "text")
```

We tokenize the model, where punct and symbols are removed.

```
In [0]: data_toks <- tokens(data_coor, what = "word") %>%  
  tokens_tolower() %>%  
  tokens(remove_punct = TRUE,  
    remove_symbols = TRUE)
```

```
In [10]: feats <- dfm(data_coor, verbose = TRUE) %>%  
  dfm_trim(min_termfreq = 5) %>%  
  featnames()  
  
data_fcm <- fcm(data_coor,  
  context = "window",  
  count = "weighted",  
  weights = 1 / (1:5),  
  tri = TRUE)
```

Creating a dfm from a corpus input...

- ... lowercasing
- ... found 18,773 documents, 159,504 features
- ... created a 18,773 x 159,504 sparse dfm
- ... complete.

Elapsed time: 8.09 seconds.

When we create the document - feature matrix and a feature-co-occurrence matrix we can use it to train GloVe embeddings.

```
In [11]: glove <- GlobalVectors$new(word_vectors_size = 50, vocabulary = featnames(data_fcm), x_max = 10)
shakes_wv_main = glove$fit_transform(data_fcm, n_iter = 10, convergence_tol = 0.001)
```

```
INFO [2019-11-29 20:19:04] 2019-11-29 20:19:04 - epoch 1, expected cost 0.1061
INFO [2019-11-29 20:19:08] 2019-11-29 20:19:08 - epoch 2, expected cost 0.0750
INFO [2019-11-29 20:19:12] 2019-11-29 20:19:12 - epoch 3, expected cost 0.0636
INFO [2019-11-29 20:19:17] 2019-11-29 20:19:17 - epoch 4, expected cost 0.0565
INFO [2019-11-29 20:19:21] 2019-11-29 20:19:21 - epoch 5, expected cost 0.0516
INFO [2019-11-29 20:19:25] 2019-11-29 20:19:25 - epoch 6, expected cost 0.0480
INFO [2019-11-29 20:19:29] 2019-11-29 20:19:29 - epoch 7, expected cost 0.0452
INFO [2019-11-29 20:19:34] 2019-11-29 20:19:34 - epoch 8, expected cost 0.0429
INFO [2019-11-29 20:19:38] 2019-11-29 20:19:38 - epoch 9, expected cost 0.0410
INFO [2019-11-29 20:19:42] 2019-11-29 20:19:42 - epoch 10, expected cost 0.0393
```

Now we create the word vectors based on the GloVe model. We want to see which words are correlated to the different topics. The tree topics we look at are 'baseball', 'christianity' and 'atheisms'.

```
In [0]: data_topi = glove$components
topi_word_vectors = shakes_wv_main + t(data_topi)
```

```
In [13]: baseball = topi_word_vectors["baseball", , drop = F]
atheism = topi_word_vectors["atheism", , drop = F]

# ham = shakes_word_vectors["hamlet", , drop = F]
cos_sim_bas = sim2(x = topi_word_vectors, y = baseball, method = "cosine", norm = "l2")
cos_sim_ath = sim2(x = topi_word_vectors, y = atheism, method = "cosine", norm = "l2")

# head(sort(cos_sim_rom[,1], decreasing = T), 10)
head(sort(cos_sim_bas[,1], decreasing = T), 5)
head(sort(cos_sim_ath [,1],decreasing = T),5)
```

baseball	1
hockey	0.850228735725524
players	0.788695876617732
team	0.746814678143087
game	0.731861223445283

atheism	1
weak	0.650023959059722
strong	0.627936921216048
faith	0.597056542506727
Its	0.593246026010867

The 5 most relevant words with the greatest correlation can be seen here. For baseball, it is seen that the word 'hokey' has the greatest correlation and then 'players'. For the subject of Atheism, the word 'strong' is the word with the highest relation. It is coded so that the top 5 words with the greatest correlation in a decreasing approach are displayed. This makes some sense in relation to the topic and what words are correlated with it.

```
In [14]: data_word_vectors <- fit_transform(data_fcm, glove, n_iter = 10)
```

```
INFO [2019-11-29 20:19:48] 2019-11-29 20:19:48 - epoch 1, expected cost 0.046
0
INFO [2019-11-29 20:19:53] 2019-11-29 20:19:53 - epoch 2, expected cost 0.039
4
INFO [2019-11-29 20:19:57] 2019-11-29 20:19:57 - epoch 3, expected cost 0.036
9
INFO [2019-11-29 20:20:01] 2019-11-29 20:20:01 - epoch 4, expected cost 0.035
2
INFO [2019-11-29 20:20:05] 2019-11-29 20:20:05 - epoch 5, expected cost 0.033
8
INFO [2019-11-29 20:20:10] 2019-11-29 20:20:10 - epoch 6, expected cost 0.032
7
INFO [2019-11-29 20:20:14] 2019-11-29 20:20:14 - epoch 7, expected cost 0.031
8
INFO [2019-11-29 20:20:18] 2019-11-29 20:20:18 - epoch 8, expected cost 0.031
0
INFO [2019-11-29 20:20:22] 2019-11-29 20:20:22 - epoch 9, expected cost 0.030
3
INFO [2019-11-29 20:20:27] 2019-11-29 20:20:27 - epoch 10, expected cost 0.02
97
```

```
In [0]: data_word_vectors %<>% as.data.frame() %>%
  rownames_to_column(var = "word") %>%
  as_tibble()

#out of this corpus we now create a tidy represation.
data_tidy2 <- data_toks %>%
  dfm() %>%
  tidy()

#and join them with the word vectors.
data_vectors <- data_tidy2 %>%
  inner_join(data_word_vectors, by = c("term" = "word"))

data_vectors %<>%
  select(-term, -count) %>%
  group_by(document) %>%
  summarise_all(mean)
```




```
In [16]: data_vectors_umap <- umap(data_vectors %>% column_to_rownames("document"),
                                   n_neighbors = 15,
                                   metric = "cosine",
                                   min_dist = 0.01,
                                   scale = TRUE,
                                   verbose = TRUE) %>%
as.data.frame()
```

```
20:20:42 UMAP embedding parameters a = 1.896 b = 0.8006
20:20:42 Read 18773 rows and found 50 numeric columns
20:20:42 Scaling to zero mean and unit variance
20:20:42 Kept 50 non-zero-variance columns
20:20:42 Using Annoy for neighbor search, n_neighbors = 15
20:20:43 Building Annoy index with metric = cosine, n_trees = 50
0%   10   20   30   40   50   60   70   80   90  100%
[----|----|----|----|----|----|----|----|----|----|
*****|
20:20:47 Writing NN index file to temp file /tmp/Rtmp02b0Xv/file8032262194
20:20:47 Searching Annoy index using 1 thread, search_k = 1500
20:20:52 Annoy recall = 100%
20:20:53 Commencing smooth kNN distance calibration using 1 thread
20:20:53 Initializing from normalized Laplacian + noise
20:20:54 Commencing optimization for 200 epochs, with 441142 positive edges
20:21:11 Optimization finished
```

We use the packages uwot to make the visualization.

```
In [17]: data_vectors_umap %>%  
  ggplot(aes(x = V1, y = V2)) +  
  geom_point(shape = 21, alpha = 0.25)
```



NETWORK

Network analysis is another way to gain insight into how the text is structured. We have previously used `unnest_tokens` as individual words, we now use `n_grams` which symbolize the order of words. We take a look at how often the word X is followed by the word Y and then build a model to see the relationship between them. When we put `n = 2` we examine pairs of each other which are also called "Bigrams"

We create the bigrams by `unnest_tokens`.

```
In [0]: bigramss <- data %>%
  unnest_tokens(bigram, text, token = "ngrams", n=2)
```

```
In [19]: bigramss$bigram[1:2]
bigramss %>% head()
```

```
'from csfed' 'csfed uxa.ecn.bgu.edu'
```

A tibble: 6 × 4

X1	text_id	topic	bigram
<dbl>	<dbl>	<chr>	<chr>
1	1366	atheism	from csfed
1	1366	atheism	csfed uxa.ecn.bgu.edu
1	1366	atheism	uxa.ecn.bgu.edu frank
1	1366	atheism	frank doss
1	1366	atheism	doss subject
1	1366	atheism	subject re

As we can see, the words overlap. The first token is "from csfed" and the second is "csfed uxa.ecn.bgu.edu". In the next step are the most common bigrams.

```
In [20]: bigramss %>%
  count(bigram, sort = TRUE) %>% head()
```

A tibble: 6 × 2

bigram	n
<chr>	<int>
of the	22150
in the	15811
subject re	12423
in article	10002
ax ax	8806
to the	8806

The most common bigrams are words like "of the", "in the" and so on. These are very common and disinterested words for analysis. These words are called "stop-words". We use now 'separate' from the packages tidyrr to split the column into two columns, "word" and "word2" and remove them if either is a stop-word.

```
In [21]: #
bigrams_separated <- bigramss %>%
  separate(bigram, c("word1", "word2"), sep = " ")

bigrams_filtered <- bigrams_separated %>%
  filter(!word1 %in% stop_words$word) %>%
  filter(!word2 %in% stop_words$word)

# new bigram counts:
bigram_counts <- bigrams_filtered %>%
  count(word1, word2, sort = TRUE)

bigram_counts %>% head()
```

A tibble: 6 × 3

word1	word2	n
<chr>	<chr>	<int>
ax	ax	8806
nntp	posting	8055
posting	host	8042
organization	university	2831
distribution	world	1687
usa	lines	1320

When the stop words are removed we now see that the common bigrams are now "ax ax", "posting host", "organization university" etc. We now want to combine two columns into a single column and this is done using the unite function from the packages tidyr. This function is the reverse of what we used earlier, 'separate'. Once we have seen how the words correlate with each other, it is interesting to visualize the relationship between the words.

```
In [22]: bigrams_united <- bigrams_filtered %>%
  unite(bigram, word1, word2, sep = " ")

bigrams_united %>% head()
```

A tibble: 6 × 4

X1	text_id	topic	bigram
<dbl>	<dbl>	<chr>	<chr>
1	1366	atheism	csfed uxa.ecn.bgu.edu
1	1366	atheism	uxa.ecn.bgu.edu frank
1	1366	atheism	frank doss
1	1366	atheism	doss subject
1	1366	atheism	theories organization
1	1366	atheism	organization educational

The words are being arranged in a network. A graph can be constructed from a tidy object since it has three variables. From are the column where the node an edge is coming from, 'to' are the node and edge is going towards and 'weight' is a numeric value associated with each edge. The alpha value is plotted as the weight, so if the bigrams occur often then the colour would be darker.

The plot shows how the words have a connection in the newsgroup. The arrows from one word to another indicate that there is a connection. As the words are lemmatized then it will not be grammatically correct, but it still makes sense to look at the connection. There are big and small connections, for example, are law -> enforcement, computer -> science and white -> house small connections, where san -> Diego -> francisco -> jose is a larger connection with several words combined to each other. At the same time, a much larger cluster is seen around the word 'lines' which is the keyword for all the numbers pointing towards each other. This makes sense since numbers are similar and therefore they are connected to each other.

```
In [24]: bigram_tf_idf <- bigrams_united %>%
  count(topic, bigram) %>%
  bind_tf_idf(bigram, topic, n) %>%
  arrange(desc(tf_idf))

bigram_tf_idf %>% head()
```

A tibble: 6 × 6

topic	bigram	n	tf	idf	tf_idf
<chr>	<chr>	<int>	<dbl>	<dbl>	<dbl>
comp-windows	ax ax	8806	0.113189116	2.995732	0.33908429
comp-windows	max ax	680	0.008740472	2.995732	0.02618411
comp-windows	ax max	678	0.008714765	2.995732	0.02610710
mideast	serdar argic	372	0.004354238	2.995732	0.01304413
crypt	clipper chip	450	0.006442653	1.897120	0.01222249
comp-windows	g9v g9v	281	0.003611872	2.995732	0.01082020

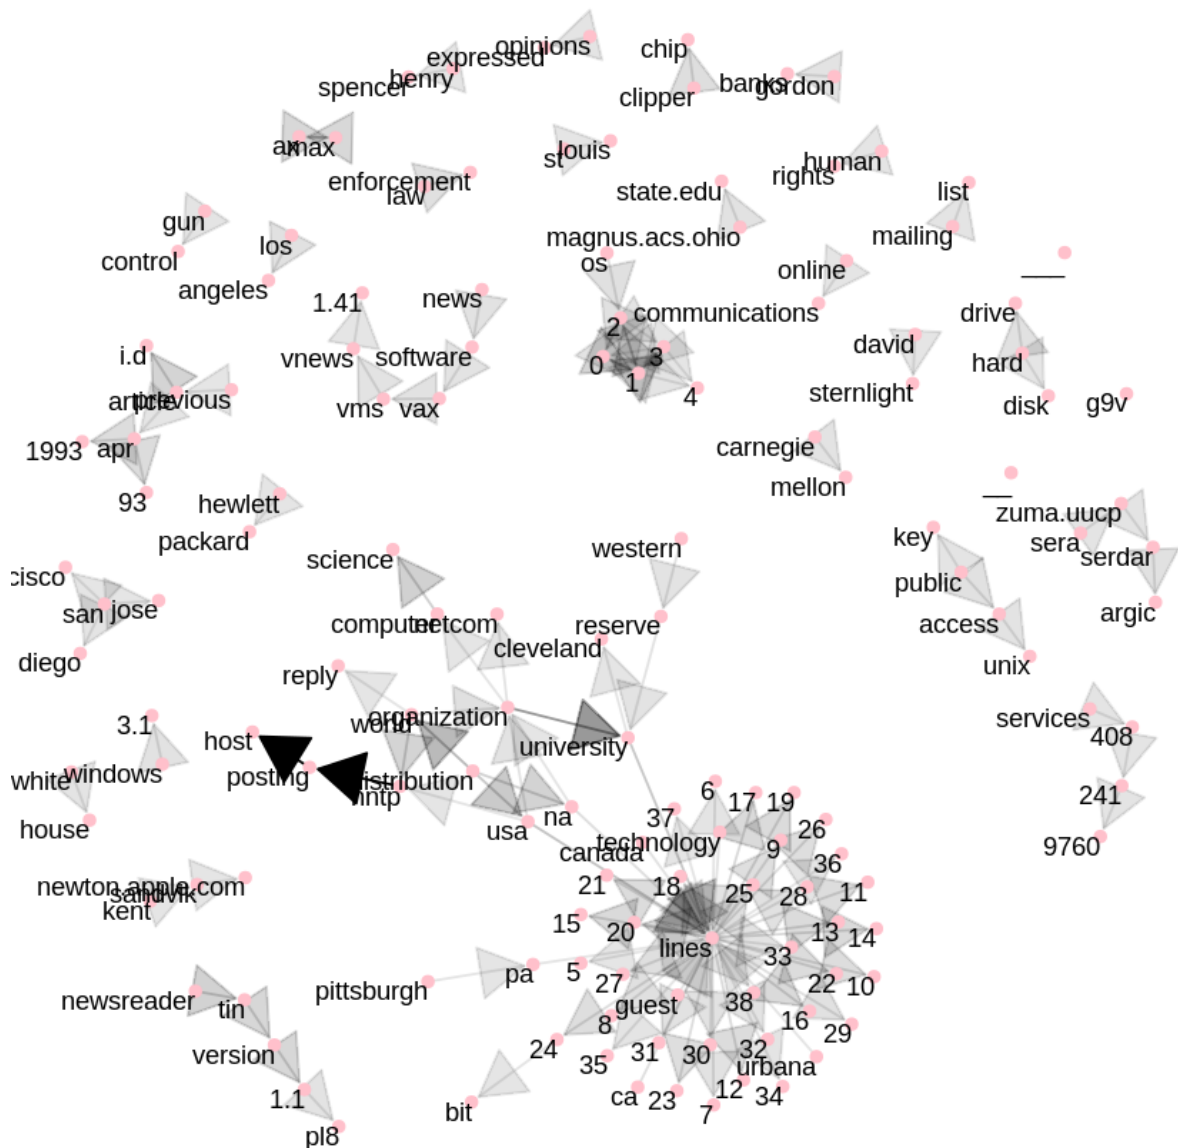
Visualizing

=====

```
In [25]: set.seed(123)
bigram_graph <- bigram_counts %>%
  filter(n > 200) %>% #The occurrence of the bigram is more than 15.
  graph_from_data_frame()

a<- grid::arrow(type = "closed",length = unit(.30,"inches"))

ggraph(bigram_graph, layout = "fr") +
  geom_edge_link(aes(edge_alpha = n), show.legend = FALSE,
    arrow = a, end_cap = circle(.05,'inches'))+
  geom_node_point(color = "pink", size = 2) +
  geom_node_text(aes(label = name),vjust=1,hjust=1) +
  theme_void()
```



We import the data

```
In [0]: import pandas as pd
import re
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import string
import nltk
import warnings

from numpy.random import seed
seed(1)

from tensorflow import set_random_seed
set_random_seed(2)
```

```
In [0]: data = pd.read_csv('https://www.dropbox.com/s/yvgkkt96fy0dp9k/File%20Name.csv?
dl=1', low_memory=False)
```

```
In [0]: df = data
```

Benchmark model

As a benchmark model we will be using simple NLP models, we will be using both a random forest model and a logistic classifier model.

This model will be used to hold our neural network up against, to see how it performs. To do this we have to make a separate training and test split, run a tf-idf on both and then run it into the models.

Random Forest

```
In [0]: from sklearn.feature_extraction.text import CountVectorizer
bow_vectorizer = CountVectorizer(max_df=0.90, min_df=2, max_features=1000, sto
p_words='english')
# bag-of-words feature matrix
bow = bow_vectorizer.fit_transform(df['text'])
```

```
In [0]: from sklearn.feature_extraction.text import TfidfVectorizer
tfidf_vectorizer = TfidfVectorizer(max_df=0.90, min_df=2, max_features=1000, s
top_words='english')
# TF-IDF feature matrix
tfidf = tfidf_vectorizer.fit_transform(df['text'])
train_bow = bow[:31962,:]
```

```
In [0]: from sklearn.model_selection import cross_val_score
        from sklearn.model_selection import train_test_split
        from sklearn.linear_model import LogisticRegression
```

```
In [0]: X_train, X_test, y_train, y_test = train_test_split(train_bow, data['topic'],
        random_state=42, test_size=0.3)
```

```
In [20]: from sklearn.ensemble import RandomForestClassifier

         RFmodel = RandomForestClassifier()

         scores = cross_val_score(RFmodel, X_train, y_train, cv = 5)
         print(scores)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: Future
Warning: The default value of n_estimators will change from 10 in version 0.2
0 to 100 in 0.22.
```

```
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: Future
Warning: The default value of n_estimators will change from 10 in version 0.2
0 to 100 in 0.22.
```

```
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: Future
Warning: The default value of n_estimators will change from 10 in version 0.2
0 to 100 in 0.22.
```

```
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: Future
Warning: The default value of n_estimators will change from 10 in version 0.2
0 to 100 in 0.22.
```

```
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: Future
Warning: The default value of n_estimators will change from 10 in version 0.2
0 to 100 in 0.22.
```

```
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

```
[0.66034156 0.65261959 0.65766451 0.66984369 0.66870229]
```

```
In [21]: from sklearn.metrics import classification_report
RFmodel.fit(X_train, y_train)

y_pred_RF = RFmodel.predict(X_test)

print(classification_report(y_test, y_pred_RF))
```

/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.

"10 in version 0.20 to 100 in 0.22.", FutureWarning)

	precision	recall	f1-score	support
atheism	0.62	0.67	0.65	225
autos	0.59	0.74	0.66	277
baseball	0.58	0.73	0.65	280
christianity	0.68	0.85	0.75	304
comp-windows	0.61	0.72	0.66	307
compgraphics	0.44	0.53	0.48	276
crypt	0.77	0.79	0.78	270
electronics	0.44	0.43	0.43	287
guns	0.69	0.76	0.72	275
hocey	0.81	0.82	0.82	319
ibm-pc-hardware	0.53	0.44	0.48	309
mac-hardware	0.68	0.60	0.64	305
medicin	0.65	0.61	0.63	330
mideast	0.89	0.83	0.86	285
misc-forsale	0.76	0.74	0.75	293
motorcycles	0.84	0.75	0.79	298
politics	0.74	0.54	0.62	231
religion-misc	0.66	0.34	0.45	188
space	0.76	0.70	0.73	289
windows-x	0.67	0.61	0.63	284
accuracy			0.67	5632
macro avg	0.67	0.66	0.66	5632
weighted avg	0.67	0.67	0.66	5632

The random forest model gives some nice results, getting a 67% accuracy with it having the most difficulty predicting the categories to do with electronics and the easiest predicting the sports related categories.

Logistic Regression

Now we will go on to a logistic regression, where the random forest relies on a series of regression trees for prediction, the logistic regression relies on linear estimations for predictions.

```
In [0]: modelLogisticRegression = LogisticRegression()
```

```
In [23]: #n-fold cross-validation (splitting the 80% into 5 part, using 4 to train and
          1 to evaluate)
          scores = cross_val_score(modelLogisticRegression, X_train, y_train, cv = 5)
          print(scores)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/logistic.py:432:
FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a s
olver to silence this warning.
```

```
FutureWarning)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/logistic.py:469:
FutureWarning: Default multi_class will be changed to 'auto' in 0.22. Specify
the multi_class option to silence this warning.
```

```
"this warning.", FutureWarning)
```

```
[0.72220114 0.72930904 0.72879422 0.7384674 0.74732824]
```

```
In [24]: modelLogisticRegression = LogisticRegression()
          modelLogisticRegression.fit(X_train, y_train)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/logistic.py:432:
FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a s
olver to silence this warning.
```

```
FutureWarning)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/logistic.py:469:
FutureWarning: Default multi_class will be changed to 'auto' in 0.22. Specify
the multi_class option to silence this warning.
```

```
"this warning.", FutureWarning)
```

```
Out[24]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, l1_ratio=None, max_iter=100,
                             multi_class='warn', n_jobs=None, penalty='l2',
                             random_state=None, solver='warn', tol=0.0001, verbose=0,
                             warm_start=False)
```

```
In [25]: print(modelLogisticRegression.score(X_test, y_test))
```

```
0.7356178977272727
```

```
In [26]: from sklearn.metrics import classification_report
y_pred_log = modelLogisticRegression.predict(X_test)

print(classification_report(y_test, y_pred_log))
```


	precision	recall	f1-score	support
atheism	0.78	0.72	0.75	225
autos	0.70	0.77	0.73	277
baseball	0.77	0.81	0.79	280
christianity	0.81	0.81	0.81	304
comp-windows	0.70	0.73	0.71	307
compgraphics	0.55	0.57	0.56	276
crypt	0.83	0.82	0.83	270
electronics	0.56	0.62	0.59	287
guns	0.77	0.79	0.78	275
hocey	0.87	0.87	0.87	319
ibm-pc-hardware	0.60	0.56	0.58	309
mac-hardware	0.71	0.68	0.69	305
medicin	0.72	0.73	0.73	330
mideast	0.90	0.90	0.90	285
misc-forsale	0.78	0.78	0.78	293
motorcycles	0.82	0.78	0.80	298
politics	0.74	0.67	0.70	231
religion-misc	0.67	0.61	0.64	188
space	0.77	0.77	0.77	289
windows-x	0.66	0.66	0.66	284
accuracy			0.74	5632
macro avg	0.74	0.73	0.73	5632
weighted avg	0.74	0.74	0.74	5632

The logistic model seems to be performing better than the random forest, we get a 74% accuracy with the model performing better in most categories but its strong and weak categories are the same as in the random forest.

LSTM

Long short-term memory (LSTM) is an artificial recurrent neural network (RNN) architecture used in the field of deep learning. Unlike standard feedforward neural networks, LSTM has feedback connections. It can not only process single data points (such as images), but also entire sequences of data (such as speech or video). For example, LSTM is commonly used for activities such as unsegmented, connected handwriting recognition, speech recognition and anomaly detection in network traffic or IDS's (intrusion detection systems).

In an LSTM network, three gates are present:



1. Input gate — This gate discovers which value from input should be used to modify the memory. Sigmoid function decides which values to let through 0,1. and tanh function gives weightage to the values which are passed deciding their level of importance ranging from -1 to 1.
2. Forget gate — discover what details to be discarded from the block. It is decided by the sigmoid function. it looks at the previous state(h_{t-1}) and the content input(x_t) and outputs a number between 0(omit this)and 1(keep this)for each number in the cell state C_{t-1} .
3. Output gate — the input and the memory of the block is used to decide the output. Sigmoid function decides which values to let through 0,1. and tanh function gives weightage to the values which are passed deciding their level of importance ranging from -1 to 1 and multiplied with output of Sigmoid.

So to summarise the cell remembers values over arbitrary time intervals and the three gates regulate the flow of information into and out of the cell. LSTM models have been used in many different areas of machine learning.

For example in 2018 OpenAI developed bots able to humans in the game of Dota 2. OpenAI Five consists of five independent but coordinated neural networks. Each network is trained by a policy gradient method without supervising teacher and contains a single-layer, 1024-unit Long-Short-Term-Memory that sees the current game

state and emits actions through several possible action heads. Also in 2019 Deepmind made a program called "AlphaStar" that used a deep LSTM core to play the complex game "Starcraft."

We import the data

```
In [4]: from keras.preprocessing.text import Tokenizer
        from keras.preprocessing.sequence import pad_sequences
```

Using TensorFlow backend.

We start by tokenizing our data.

- We start by tokenizing our data.
- Vectorize tweet text, by turning each text into either a sequence of integers or into a vector.
- Limit the data set to the top 50.000
- Set the max number of words to be 250.

```
In [5]: # The maximum number of words to be used. (most frequent)
        MAX_NB_WORDS = 50000
        # Max number of words in each complaint.
        MAX_SEQUENCE_LENGTH = 250
        # This is fixed.
        EMBEDDING_DIM = 100
        tokenizer = Tokenizer(num_words=MAX_NB_WORDS, filters='!"#$%&()*+,-./:;<=>?@
        [\]^_`{|}~', lower=True)
        tokenizer.fit_on_texts(df['text'].values)
        word_index = tokenizer.word_index
        print('Found %s unique tokens.' % len(word_index))
```

Found 164987 unique tokens.

Next step is to truncate and pad the input sequences so that they are all in the same length for modeling.

```
In [6]: X = tokenizer.texts_to_sequences(df['text'].values)
        X = pad_sequences(X, maxlen=MAX_SEQUENCE_LENGTH)
        print('Shape of data tensor:', X.shape)
```

Shape of data tensor: (18773, 250)

```
In [7]: Y = pd.get_dummies(df['topic']).values
        print('Shape of label tensor:', Y.shape)
```

Shape of label tensor: (18773, 20)

After tokenizing our data it's time to split our data into train and test.

```
In [0]: from sklearn.model_selection import train_test_split
```

```
In [9]: X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size = 0.10, ran
dom_state = 42)
print(X_train.shape,Y_train.shape)
print(X_test.shape,Y_test.shape)

(16895, 250) (16895, 20)
(1878, 250) (1878, 20)
```

It is now finally time to create our LSTM model.

```
In [0]: from keras.models import Sequential
from keras.layers import Embedding
from keras.layers import SpatialDropout1D
from keras.layers import LSTM
from keras.layers import Dense
from keras.callbacks import ModelCheckpoint, EarlyStopping
```

```
In [0]: import keras
```

The LSTM model in this project is constructed as follows:

- The first layer is the embedded layer that uses 100 length vectors to represent each word.
- SpatialDropout1D performs variational dropout in NLP models.
- The next layer is the LSTM layer with 100 memory units.
- The output layer must create 20 output values, one for each outcome
- Activation function is softmax for multi-class classification.
- Because it is a multi-class classification problem, categorical_crossentropy is used as the loss function.
- 10 epochs have been used and gives the most accurate model without taking too much loading time. A more accurate model is perhaps possible with more epochs but it would increase the loading time significantly so this has been excluded and maintained at 10 epoch.

Various other inputs have been tested, but constructing our model with these specific values has constructed the most accurate model, which the reader is free to test for themselves.


```
In [12]: model = Sequential()
model.add(Embedding(MAX_NB_WORDS, EMBEDDING_DIM, input_length=X.shape[1]))
model.add(SpatialDropout1D(0.2))
model.add(LSTM(100, dropout=0.2, recurrent_dropout=0.2))
model.add(Dense(20, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

epochs = 10
batch_size = 64

history = model.fit(X_train, Y_train, epochs=epochs, batch_size=batch_size, validation_split=0.1, callbacks=[EarlyStopping(monitor='val_loss', patience=3, min_delta=0.0001)])
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:66: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4432: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:148: The name tf.placeholder_with_default is deprecated. Please use tf.compat.v1.placeholder_with_default instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3733: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3576: The name tf.log is deprecated. Please use tf.math.log instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/core/python/ops/math_grad.py:1424: where (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1033: The name tf.assign_add is deprecated. Please use tf.compat.v1.assign_add instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1020: The name tf.assign is deprecated. Please use tf.compat.v1.assign instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3005: The name tf.Session is deprecated. Please use tf.compat.v1.Session instead.

Train on 15205 samples, validate on 1690 samples

Epoch 1/10

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:190: The name tf.get_default_session is deprecated. Please use tf.compat.v1.get_default_session instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:197: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:207: The name tf.global_variables is deprecated. Please use tf.compat.v1.global_variables instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:216: The name tf.is_variable_initialized is deprecated. Please use tf.compat.v1.is_variable_initialized instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:223: The name tf.variables_initializer is deprecated. Please use tf.compat.v1.variables_initializer instead.

```
15205/15205 [=====] - 141s 9ms/step - loss: 2.7915 -
acc: 0.1237 - val_loss: 2.4287 - val_acc: 0.2207
Epoch 2/10
15205/15205 [=====] - 133s 9ms/step - loss: 2.1354 -
acc: 0.3331 - val_loss: 1.8454 - val_acc: 0.4420
Epoch 3/10
15205/15205 [=====] - 134s 9ms/step - loss: 1.4736 -
acc: 0.5240 - val_loss: 1.5684 - val_acc: 0.5041
Epoch 4/10
15205/15205 [=====] - 134s 9ms/step - loss: 1.0488 -
acc: 0.6676 - val_loss: 1.4385 - val_acc: 0.5716
Epoch 5/10
15205/15205 [=====] - 134s 9ms/step - loss: 0.7588 -
acc: 0.7628 - val_loss: 1.3721 - val_acc: 0.6030
Epoch 6/10
15205/15205 [=====] - 133s 9ms/step - loss: 0.5713 -
acc: 0.8266 - val_loss: 1.4779 - val_acc: 0.5959
Epoch 7/10
15205/15205 [=====] - 133s 9ms/step - loss: 0.4667 -
acc: 0.8606 - val_loss: 1.4281 - val_acc: 0.6343
Epoch 8/10
15205/15205 [=====] - 134s 9ms/step - loss: 0.3217 -
acc: 0.9036 - val_loss: 1.4499 - val_acc: 0.6491
```

The model has been constructed. It is now time to interpret the results

In [13]: `print(model.summary())`

Model: "sequential_1"

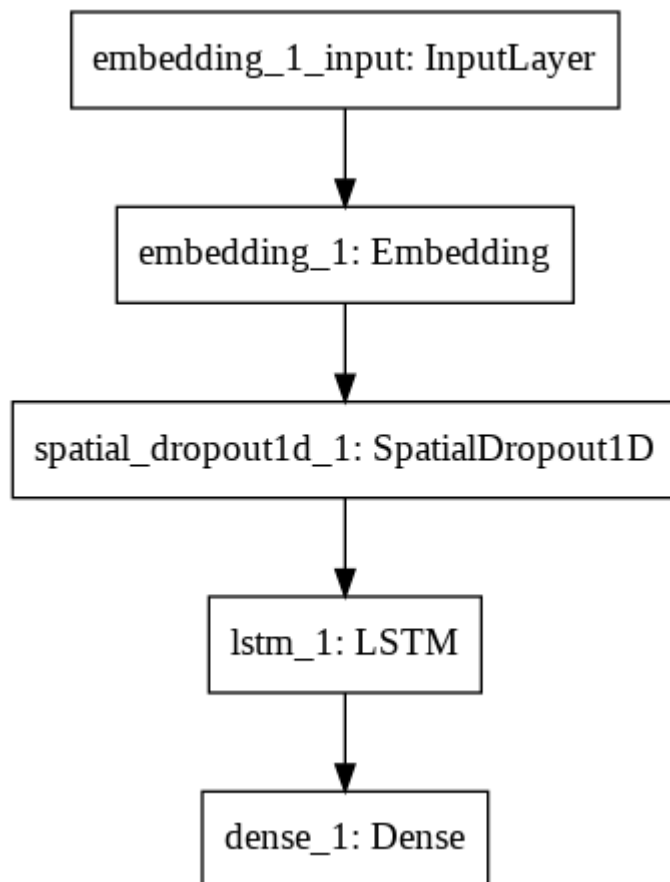
Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 250, 100)	5000000
spatial_dropout1d_1 (Spatial Dropout)	(None, 250, 100)	0
lstm_1 (LSTM)	(None, 100)	80400
dense_1 (Dense)	(None, 20)	2020
Total params: 5,082,420		
Trainable params: 5,082,420		
Non-trainable params: 0		
None		

In [0]: `from keras.utils.vis_utils import plot_model`

The Keras package has an included visualization package which helps explain to the reader how the model is constructed. This visualization will be presented below.

In [15]: `plot_model(model, to_file='model.png')`

Out[15]:



Now it's time for the true test. We are going to test the model on the actual test set.

```
In [24]: accr = model.evaluate(X_test,Y_test)
print('Test set\n Loss: {:.3f}\n Accuracy: {:.3f}'.format(accr[0],accr[1]))
```

1878/1878 [=====] - 12s 6ms/step
Test set
Loss: 1.564
Accuracy: 0.619

We get an Accuracy of 61.9% which is quite amazing considering how many categories there are and how little difference there is between them. The accuracy might seem low, but it is important to consider the model has to classify between 20 different variables which is quite problematic. Even a human would get massive errors if it had to classify a text between 20 different text topics. We will now further explore the model to see which topics are hard to classify.

```
In [0]: from sklearn.metrics import classification_report
y_pred = model.predict(X_test)
```

```
In [0]: y1 = np.argmax(y_pred, axis=1)
```

```
In [0]: y2 = np.argmax(Y_test, axis=1)
```


```
In [0]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y2,y1)
```

```
In [25]: print(classification_report(y2, y1, target_names=['Atheism', 'Autos', 'Baseball', 'Christianity', 'Compgraphics', 'Comp-windows', 'Crypt', 'Electronics', 'Guns', 'Hockey', 'IBM-PC-hardware', 'Mac-hardware', 'Medicin', 'Mideast', 'Misc-forsale', 'Motorcycles', 'Politics', 'Religion-misc', 'Space', 'Windows-x',]))
```


	precision	recall	f1-score	support
Atheism	0.64	0.74	0.69	58
Autos	0.48	0.53	0.51	83
Baseball	0.76	0.83	0.79	105
Christianity	0.60	0.69	0.64	81
Compgraphics	0.60	0.60	0.60	110
Comp-windows	0.44	0.39	0.41	90
Crypt	0.78	0.57	0.66	102
Electronics	0.41	0.37	0.39	98
Guns	0.63	0.65	0.64	88
Hockey	0.83	0.82	0.83	91
IBM-PC-hardware	0.52	0.56	0.54	109
Mac-hardware	0.58	0.57	0.58	110
Medicin	0.62	0.62	0.62	108
Mideast	0.60	0.82	0.69	87
Misc-forsale	0.72	0.60	0.65	92
Motorcycles	0.87	0.60	0.71	97
Politics	0.49	0.61	0.54	82
Religion-misc	0.46	0.43	0.44	75
Space	0.69	0.62	0.66	101
Windows-x	0.69	0.77	0.73	111
accuracy			0.62	1878
macro avg	0.62	0.62	0.62	1878
weighted avg	0.63	0.62	0.62	1878

We can observe from the above figure that the model can easier classify certain topics compared to others. For example the topic "Hockey" has a accuracy of 83%. The LSTM model experiences problems when having to classify Electronics, which can be explained through the other topics. There is so many topics related to electronics for example Windows X, MAC hardware, IBM-PC-hardware and compgraphics, which makes this topic and other topics hard to classify. IBM-PC-hardware and Mac-hardware accuracy is also quite low on 52% and 58%. Same for Comp windows on 44%, which indicates the model mixes the different topic with each other. Baseball is good example of a topic that is not related to many other topics. This is also where the model achieves very high accuracy. Same for Motorcycles and so on. Overall we can conclude the LSTM neural network is quite accurate, but more accurate for certain topics.

```
In [22]: plt.title('Accuracy')
plt.plot(history.history['acc'], label='train')
plt.plot(history.history['val_acc'], label='test')
plt.legend()
plt.show();
```



```
In [23]: plt.title('Loss')
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='test')
plt.legend()
plt.show();
```



The plots suggest that the model has a little over fitting problem, more data may help, but more epochs will not help using the current data and will only increase only loading times. More data for all topics could help reduce this over fitting problem.

Conclusion

For this conclusion we will be focusing primarily on our benchmark models and the neural network, the Glove model helps us with identifying words that are important in each topic. We find that electronics are connected as well as sport topics, and that some combinations of words are used in specific categories.

Looking at the 3 models, we see that the logistic and random forest models are outperforming our neural network. While not expected, this makes some sense as the classification models are based on a bag of word approach making the words everything and as there will be words used exclusively in each topic. The neural network takes sequences into a count, the sequence of words is often a personal thing and will therefore work well in identifying a person from writing but less in identifying a topic. Further though our model has been run for some time it still has room for improvements, both by adding epochs and more layers and units in each layer and therefore there are room for improvements.